

NIJs Recidivism

In this project, KMeans clustering was applied as an unsupervised learning approach to group individuals in the NIJ Recidivism dataset based on demographic, employment, and criminal history features. The goal of clustering was to identify latent patterns among individuals without using recidivism outcomes as labels, allowing the data itself to reveal groups with similar characteristics. This unsupervised approach is appropriate because it does not rely on pre-defined outcomes, and it can uncover meaningful profiles that might inform targeted interventions or rehabilitation strategies.

Citation

National Institute of Justice. (2021). NIJ's Recidivism Challenge Full Dataset [Dataset]. Office of Justice Programs, U.S. Department of Justice. Retrieved from <https://catalog.data.gov/dataset/nijs-recidivism-challenge-full-dataset>

Set-Up

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score

sns.set(style="whitegrid")
```

EDA

Initial data exploration revealed a mix of numeric and categorical features with varying levels of missing data. Numeric variables such as Percent_Days_Employed, Jobs_Per_Year, and Age_at_Release were cleaned by converting text ranges and categories into numeric representations and filling missing values with the median. Categorical features, including Gender and Race, were one-hot encoded. Distributional analysis showed that most individuals were male, aged between 30 and 35 at release, and racially diverse. Employment levels varied widely, from near-zero employment to full employment, while supervision risk scores were generally moderate to high. Correlation analysis indicated that employment metrics were negatively correlated with recidivism, while risk scores were slightly positively correlated. Standard scaling was applied to numeric features in preparation for PCA and clustering.

```
In [2]: df = pd.read_csv('NIJ_s_Recidivism_Challenge_Full_Dataset.csv')

df.head()
```

```
Out[2]:   ID  Gender  Race  Age_at_Release  Residence_PUMA  Gang_Affiliated  Supervision_Risk
0    1      M  BLACK       43-47             16        False
1    2      M  BLACK       33-37             16        False
2    3      M  BLACK    48 or older            24        False
3    4      M  WHITE      38-42             16        False
4    5      M  WHITE      33-37             16        False
```

5 rows × 54 columns

```
In [3]: df.info()

df.describe()

df.isnull().sum()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 25835 entries, 0 to 25834
Data columns (total 54 columns):
 #   Column           Non-Null Count Dtype
 ---  -- 
 0   ID               25835 non-null  int64
 1   Gender            25835 non-null  object
 2   Race              25835 non-null  object
 3   Age_at_Release    25835 non-null  object
 4   Residence_PUMA    25835 non-null  int64
 5   Gang_Affiliated   22668 non-null  object
 6   Supervision_Risk_Score_First  25360 non-null  float64
 7   Supervision_Level_First   24115 non-null  object
 8   Education_Level     25835 non-null  object
 9   Dependents          25835 non-null  object
 10  Prison_Offense     22558 non-null  object
 11  Prison_Years       25835 non-null  object
 12  Prior_Arrest_Episodes_Felony  25835 non-null  object
 13  Prior_Arrest_Episodes_Misd   25835 non-null  object
 14  Prior_Arrest_Episodes_Violent 25835 non-null  object
 15  Prior_Arrest_Episodes_Property 25835 non-null  object
 16  Prior_Arrest_Episodes_Drug   25835 non-null  object
 17  Prior_Arrest_Episodes_PPViolationCharges 25835 non-null  object
 18  Prior_Arrest_Episodes_DVCharges 25835 non-null  bool
 19  Prior_Arrest_Episodes_GunCharges 25835 non-null  bool
 20  Prior_Conviction_Episodes_Felony  25835 non-null  object
 21  Prior_Conviction_Episodes_Misd   25835 non-null  object
 22  Prior_Conviction_Episodes_Viol  25835 non-null  bool
 23  Prior_Conviction_Episodes_Prop  25835 non-null  object
 24  Prior_Conviction_Episodes_Drug  25835 non-null  object
 25  Prior_Conviction_Episodes_PPViolationCharges 25835 non-null  bool
 26  Prior_Conviction_Episodes_DomesticViolenceCharges 25835 non-null  bool
 27  Prior_Conviction_Episodes_GunCharges 25835 non-null  bool
 28  Prior_Revocations_Parole     25835 non-null  bool
 29  Prior_Revocations_Probation  25835 non-null  bool
 30  Condition_MH_SA          25835 non-null  bool
 31  Condition_Cog_Ed         25835 non-null  bool
 32  Condition_Other          25835 non-null  bool
 33  Violations_ElectronicMonitoring 25835 non-null  bool
 34  Violations_Instruction    25835 non-null  bool
 35  Violations_FailToReport   25835 non-null  bool
 36  Violations_MoveWithoutPermission 25835 non-null  bool
 37  Delinquency_Report       25835 non-null  object
 38  Program_Attendances     25835 non-null  object
 39  Program_UnexcusedAbsences 25835 non-null  object
 40  Residence_Changes       25835 non-null  object
 41  Avg_Days_per_DrugTest   19732 non-null  float64
 42  DrugTests_THC_Positive  20663 non-null  float64
 43  DrugTests_Cocaine_Positive 20663 non-null  float64
 44  DrugTests_Meth_Positive  20663 non-null  float64
 45  DrugTests_Other_Positive 20663 non-null  float64
 46  Percent_Days_Employed   25373 non-null  float64
 47  Jobs_Per_Year          25027 non-null  float64
 48  Employment_Exempt      25835 non-null  bool
 49  Recidivism_Within_3years 25835 non-null  bool
 50  Recidivism_Arrest_Year1  25835 non-null  bool
```

```
51 Recidivism_Arrest_Year2           25835 non-null  bool
52 Recidivism_Arrest_Year3           25835 non-null  bool
53 Training_Sample                  25835 non-null  int64
dtypes: bool(20), float64(8), int64(3), object(23)
memory usage: 7.2+ MB
```

```
Out[3]: ID          0
Gender        0
Race          0
Age_at_Release 0
Residence_PUMA 0
Gang_Affiliated 3167
Supervision_Risk_Score_First 475
Supervision_Level_First 1720
Education_Level 0
Dependents     0
Prison_Offense 3277
Prison_Years    0
Prior_Arrest_Episodes_Felony 0
Prior_Arrest_Episodes_Misd 0
Prior_Arrest_Episodes_Violent 0
Prior_Arrest_Episodes_Property 0
Prior_Arrest_Episodes_Drug 0
Prior_Arrest_Episodes_PPViolationCharges 0
Prior_Arrest_Episodes_DVCharges 0
Prior_Arrest_Episodes_GunCharges 0
Prior_Conviction_Episodes_Felony 0
Prior_Conviction_Episodes_Misd 0
Prior_Conviction_Episodes_Viol 0
Prior_Conviction_Episodes_Prop 0
Prior_Conviction_Episodes_Drug 0
Prior_Conviction_Episodes_PPViolationCharges 0
Prior_Conviction_Episodes_DomesticViolenceCharges 0
Prior_Conviction_Episodes_GunCharges 0
Prior_Revocations_Parole 0
Prior_Revocations_Probation 0
Condition_MH_SA 0
Condition_Cog_Ed 0
Condition_Other 0
Violations_ElectronicMonitoring 0
Violations_Instruction 0
Violations_FailToReport 0
Violations_MoveWithoutPermission 0
Delinquency_Report 0
Program_Attendances 0
Program_UnexcusedAbsences 0
Residence_Changes 0
Avg_Days_per_DrugTest 6103
DrugTests_THC_Positive 5172
DrugTests_Cocaine_Positive 5172
DrugTests_Meth_Positive 5172
DrugTests_Other_Positive 5172
Percent_Days_Employed 462
Jobs_Per_Year 808
Employment_Exempt 0
Recidivism_Within_3years 0
Recidivism_Arrest_Year1 0
Recidivism_Arrest_Year2 0
Recidivism_Arrest_Year3 0
Training_Sample 0
dtype: int64
```

```
In [4]: # unique values in these columns
for col in ['Dependents', 'Jobs_Per_Year']:
    print(col, df[col].unique())
def convert_to_numeric(x):
    if pd.isnull(x):
        return np.nan
    elif isinstance(x, str) and 'or more' in x:
        return float(x.split()[0]) # take the first number
    else:
        return float(x)

df['Dependents_numeric'] = df['Dependents'].apply(convert_to_numeric)
df['Jobs_Per_Year_numeric'] = df['Jobs_Per_Year'].apply(convert_to_numeric)
```

```
Dependents ['3 or more' '1' '0' '2']
Jobs_Per_Year [0.44761029 2.           0.           ... 0.2818287 1.37484316 1.3346528
6]
```

```
In [5]: # convert ranges to midpoint
def range_to_midpoint(x):
    if isinstance(x, str) and '-' in x:
        start, end = x.split('-')
        return (float(start) + float(end)) / 2
    else:
        return np.nan

df['Age_at_Release_numeric'] = df['Age_at_Release'].apply(range_to_midpoint)
```

Visualize

Exploratory visualization confirmed these patterns: Histograms of age, supervision risk scores, and employment showed diverse distributions, while countplots revealed gender and race breakdowns. A PCA projection provided a preliminary view of feature relationships in two-dimensional space, facilitating the subsequent clustering analysis.


```
In [6]: # Age distribution
sns.histplot(df['Age_at_Release_numeric'], bins=20)
plt.xlabel("Age at Release")
plt.title("Age at Release Distribution")
plt.show()



# Gender distribution
sns.countplot(x='Gender', data=df)
plt.title("Gender Distribution")
plt.show()


# Race distribution
sns.countplot(x='Race', data=df)
plt.title("Race Distribution")
plt.show()

# Supervision risk score
```

```
sns.histplot(df['Supervision_Risk_Score_First'], bins=20)
plt.title("Supervision Risk Score Distribution")
plt.show()
```








```
In [7]: # Replace numeric_features
numeric_features = [
    'Age_at_Release_numeric',
    'Supervision_Risk_Score_First',
    'Dependents_numeric',
    'Percent_Days_Employed',
    'Jobs_Per_Year_numeric'
]

# Fill missing values
df[numeric_features] = df[numeric_features].fillna(df[numeric_features].median())
```

```
In [8]: # Correlation heatmap
corr = df[numeric_features].corr()
sns.heatmap(corr, annot=True, cmap='coolwarm')
plt.title("Correlation Heatmap")
plt.show()
```




```
In [9]: total_count = len(df)

# Count individuals with prior violent convictions
violent_count = df['Prior_Conviction_Episodes_Viol'].sum()

labels = ['Total Population', 'Violent Crime']
counts = [total_count, violent_count]

# Plot
plt.figure(figsize=(6,4))
plt.bar(labels, counts, color=['skyblue', 'salmon'])
plt.ylabel('Number of Individuals')
plt.title('Violent Crime vs Total Population')
plt.show()
```



```
In [10]: # Fill missing numeric values with median
df[numeric_features] = df[numeric_features].fillna(df[numeric_features].median())


categorical_features = ['Race', 'Gender']
df_encoded = pd.get_dummies(df[categorical_features], drop_first=True)

# Combine numeric + encoded categorical features
X = pd.concat([df[numeric_features], df_encoded], axis=1)

# Scale data for clustering/PCA
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

```
In [11]: pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_scaled)

plt.figure(figsize=(8,6))
plt.scatter(X_pca[:,0], X_pca[:,1], alpha=0.5)
plt.xlabel('PC1')
plt.ylabel('PC2')
plt.title('PCA Projection of NIJ Dataset')
plt.show()
```



KMeans

KMeans clustering was performed on scaled numeric and encoded categorical features, with the optimal number of clusters determined via silhouette analysis as four. Clusters were profiled based on age, risk score, dependents, employment, and jobs per year.

The results revealed distinct patterns: one cluster with very low employment exhibited the highest recidivism rate (71%), whereas clusters with higher employment had lower recidivism rates (45–53%). PCA visualization confirmed that the clusters were separable in two-dimensional space. Cluster labeling based on employment and risk levels provided interpretable insights, suggesting that employment is a critical protective factor against recidivism and that low-employment individuals represent a high-risk group.


```
In [12]: # Determine optimal number of clusters
sil_scores = []
K_range = range(2, 11)
for k in K_range:
    kmeans = KMeans(n_clusters=k, random_state=42)
    labels = kmeans.fit_predict(X_scaled)
    sil_scores.append(silhouette_score(X_scaled, labels))

# silhouette scores
```

```
plt.plot(K_range, sil_scores, marker='o')
plt.xlabel('Number of clusters')
plt.ylabel('Silhouette Score')
plt.title('KMeans Silhouette Analysis')
plt.show()
```


C:\Users\asm3886\anaconda3\Lib\site-packages\threadpoolctl.py:1226: RuntimeWarning:
Found Intel OpenMP ('libiomp') and LLVM OpenMP ('libomp') loaded at
the same time. Both libraries are known to be incompatible and this
can cause random crashes or deadlocks on Linux when loaded in the
same Python program.
Using threadpoolctl may cause crashes or deadlocks. For more
information and possible workarounds, please see
https://github.com/joblib/threadpoolctl/blob/master/multiple_openmp.md

```
warnings.warn(msg, RuntimeWarning)
```



```
In [13]: # Fit KMeans
kmeans = KMeans(n_clusters=4, random_state=42)
clusters = kmeans.fit_predict(X_scaled)
df['cluster'] = clusters

# Cluster in PCA space
plt.figure(figsize=(8,6))
sns.scatterplot(x=X_pca[:,0], y=X_pca[:,1], hue=df['cluster'], palette='Set2')
plt.title('Clusters Visualized on PCA Components')
plt.show()
```



```
In [14]: # mean of numeric features
cluster_summary = df.groupby('cluster')[numeric_features].mean()

# mean recidivism rate
recidivism_rate = df.groupby('cluster')['Recidivism_Within_3years'].mean()

# Combine summaries
cluster_info = cluster_summary.copy()
cluster_info['Recidivism_Rate'] = recidivism_rate

print(cluster_info)
```

```
Age_at_Release_numeric  Supervision_Risk_Score_First  \
cluster
0                      31.712495                  6.012110
1                      32.792531                  5.877510
2                      30.568253                  6.263661
3                      33.074250                  6.063191

Dependents_numeric  Percent_Days_Employed  Jobs_Per_Year_numeric  \
cluster
0                  1.647386          0.807555          1.182983
1                  1.335934          0.794549          1.242640
2                  1.398816          0.045885          0.167190
3                  1.558294          0.493171          0.747794

Recidivism_Rate
cluster
0          0.488389
1          0.527303
2          0.712445
3          0.455292
```

```
In [15]: def name_cluster(row):
    # Risk Level
    if row['Recidivism_Rate'] > 0.50:
        risk = "High-Risk"
    elif row['Recidivism_Rate'] > 0.2:
        risk = "Moderate-Risk"
    else:
        risk = "Lower-Risk"

    # Employment descriptor
    emp = row['Percent_Days_Employed']
    if emp > 0.7:
        emp_label = "High Employment"
    elif emp > 0.4:
        emp_label = "Partial Employment"
    else:
        emp_label = "Low Employment"

    return f"{risk} / {emp_label}"
```

```
In [16]: # Create a mapping from cluster number → name
cluster_names = cluster_info.apply(name_cluster, axis=1).to_dict()

# Map
df['cluster_name'] = df['cluster'].map(cluster_names)

# Preview
df[['cluster', 'cluster_name']].drop_duplicates()
```

```
Out[16]:   cluster      cluster_name
0         0  Moderate-Risk / High Employment
2         2  High-Risk / Low Employment
3         1  High-Risk / High Employment
8         3  Moderate-Risk / Partial Employment
```

```
In [17]: sns.scatterplot(
    x=X_pca[:,0],
    y=X_pca[:,1],
    hue=df['cluster_name'],
    palette='Set2',
    alpha=0.6
)
plt.title('Names')
plt.xlabel('PC1')
plt.ylabel('PC2')
plt.legend(title='Cluster')
plt.show()
```



```
In [18]: # cluster profiles
cluster_summary = df.groupby('cluster')[numeric_features].mean()
print(cluster_summary)

# Compare recidivism rates
```

```
recidivism_rate_by_cluster = df.groupby('cluster')['Recidivism_Within_3years'].mean()
print(recidivism_rate_by_cluster)

   Age_at_Release_numeric  Supervision_Risk_Score_First \
cluster
0                  31.712495          6.012110
1                  32.792531          5.877510
2                  30.568253          6.263661
3                  33.074250          6.063191

   Dependents_numeric  Percent_Days_Employed  Jobs_Per_Year_numeric \
cluster
0             1.647386          0.807555          1.182983
1             1.335934          0.794549          1.242640
2             1.398816          0.045885          0.167190
3             1.558294          0.493171          0.747794

cluster
0    0.488389
1    0.527303
2    0.712445
3    0.455292
Name: Recidivism_Within_3years, dtype: float64
```

```
In [19]: recid_cols = ['Recidivism_Arrest_Year1', 'Recidivism_Arrest_Year2', 'Recidivism_Arr


# Melt the dataframe
df_long = df.melt(
    id_vars=['ID', 'Percent_Days_Employed'], # keep ID and employment
    value_vars=recid_cols,
    var_name='Year',
    value_name='Recidivism'
)

# 'Year' to numeric 1,2,3
df_long['Year'] = df_long['Year'].str.extract(r'(\d)').astype(int)
```

```
In [20]: # Select only individuals who did NOT reoffend within 3 years
df_no_recid = df[df['Recidivism_Within_3years'] == False]

# Employment distribution
plt.figure(figsize=(8,6))
sns.histplot(df_no_recid['Percent_Days_Employed'], bins=10, kde=True)
plt.xlabel('Percent Days Employed')
plt.ylabel('Number of Individuals')
plt.title('Employment Distribution for Non-Recidivists')
plt.show()

# Categorize employment levels
df_no_recid['Employment_Level'] = pd.cut(
    df_no_recid['Percent_Days_Employed'],
    bins=[-0.01, 0.4, 0.7, 1],
    labels=['Low', 'Partial', 'High']
)
```



C:\Users\asm3886\AppData\Local\Temp\ipykernel_10708\1606978834.py:13: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy


```
df_no_recid['Employment_Level'] = pd.cut(
```

```
In [21]: df_mean = df_long.groupby(['Year', pd.cut(df_long['Percent_Days_Employed'], bins=[-1, 0, 10, 20, 30, 40, 50, 60, 70, 80, 90, 100]]), as_index=False).mean().dropna()

plt.figure(figsize=(8,6))
sns.lineplot(
    x='Year',
    y='Recidivism',
    hue='Percent_Days_Employed',
    data=df_mean,
    marker='o'
)
plt.xlabel('Years Since Release')
plt.ylabel('Average Recidivism Rate')
plt.title('Average Recidivism Over Time by Employment Level')
plt.show()
```

C:\Users\asm3886\AppData\Local\Temp\ipykernel_10708\334640896.py:1: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.


```
df_mean = df_long.groupby(['Year', pd.cut(df_long['Percent_Days_Employed'], bins=[-0.01, 0.4, 0.7, 1], labels=['Low', 'Partial', 'High'])])['Recidivism'].mean().reset_index()
```



```
In [22]: # Categorize employment levels for the whole dataset
df['Employment_Level'] = pd.cut(
    df['Percent_Days_Employed'],
    bins=[-0.01, 0.4, 0.7, 1],
    labels=['Low', 'Partial', 'High']
)

#column for recidivism status label
df['Recidivism_Status'] = df['Recidivism_Within_3years'].map({True: 'Recidivist', False: 'Non-Recidivist'})

# Plot side-by-side countplot
plt.figure(figsize=(10,6))
sns.countplot(x='Employment_Level', hue='Recidivism_Status', data=df, order=['Low', 'Partial', 'High'])
plt.xlabel('Employment Level')
plt.ylabel('Number of Individuals')
plt.title('Employment Levels: Recidivists vs Non-Recidivists')
plt.legend(title='Recidivism Status')
plt.show()
```



```
In [23]: numeric_features = [
    'Age_at_Release_numeric',
    'Supervision_Risk_Score_First',
    'Dependents_numeric',
    'Percent_Days_Employed',
    'Jobs_Per_Year_numeric'
]

# categorical features
categorical_features = [
    'Gender',
    'Race',
    'Education_Level',
    'Supervision_Level_First',
    'Gang_Affiliated',
    'Employment_Exempt'
]

# Function to get mode of categorical feature
def mode_or_na(series):
    if series.mode().empty:
        return None
    return series.mode()[0]

# Group by cluster
cluster_summary = df.groupby('cluster').agg(
    **{f'{col}_avg': (col, 'mean') for col in numeric_features},
    **{f'{col}_mode': (col, mode_or_na) for col in categorical_features},
    Recidivism_Rate=('Recidivism_Within_3years', 'mean')
).reset_index()

# Map cluster numbers to cluster names
```

```
if 'cluster_name' in df.columns:
    cluster_summary['cluster_name'] = cluster_summary['cluster'].map(
        df.set_index('cluster')['cluster_name'].to_dict()
    )

# Reorder columns
cols_order = ['cluster', 'cluster_name'] + \
    [f'{col}_avg' for col in numeric_features] + \
    [f'{col}_mode' for col in categorical_features] + \
    ['Recidivism_Rate']

cluster_summary = cluster_summary[cols_order]

# Display the summary
print(cluster_summary)
```

	cluster	cluster_name	Age_at_Release_numeric_avg	
0	0	Moderate-Risk / High Employment	31.712495	
1	1	High-Risk / High Employment	32.792531	
2	2	High-Risk / Low Employment	30.568253	
3	3	Moderate-Risk / Partial Employment	33.074250	

	Supervision_Risk_Score_First_avg	Dependents_numeric_avg	
0	6.012110	1.647386	
1	5.877510	1.335934	
2	6.263661	1.398816	
3	6.063191	1.558294	

	Percent_Days_Employed_avg	Jobs_Per_Year_numeric_avg	Gender_mode	Race_mode	
0	0.807555	1.182983	M	BLACK	
1	0.794549	1.242640	M	WHITE	
2	0.045885	0.167190	M	BLACK	
3	0.493171	0.747794	F	WHITE	

	Education_Level_mode	Supervision_Level_First_mode	Gang_Affiliated_mode	
0	High School Diploma	Standard	False	
1	High School Diploma	Standard	False	
2	Less than HS diploma	Standard	False	
3	High School Diploma	Standard	None	

	Employment_Exempt_mode	Recidivism_Rate
0	False	0.488389
1	False	0.527303
2	False	0.712445
3	False	0.455292

```
In [24]: key_features = [
    'Age_at_Release_numeric',
    'Percent_Days_Employed',
    'Jobs_Per_Year_numeric',
    'Supervision_Risk_Score_First',
    'Dependents_numeric',
    'Recidivism_Within_3years'
]

#Cluster means
```

```

cluster_means = df.groupby('cluster')[key_features].mean()


# Center
cluster_means_centered = cluster_means - df[key_features].mean()

if 'cluster_name' in df.columns:
    cluster_means_centered.index = df.groupby('cluster')['cluster_name'].first()

# Sort clusters by Age
cluster_means_centered = cluster_means_centered.sort_values('Age_at_Release_numeric')

# Plot heatmap
plt.figure(figsize=(12,6))
sns.heatmap(
    cluster_means_centered,
    annot=True,
    cmap='RdBu_r',
    center=0,
    fmt=".2f",
    linewidths=0.5,
    cbar_kws={'label': 'Deviation from Mean'}
)
plt.title("Risk and Employment vs Features")
plt.xlabel("Feature")
plt.ylabel("Cluster")
plt.xticks(rotation=45, ha='right')
plt.show()

```



```

In [25]: # Cluster means
cleaned_csv = cluster_means.copy()

if 'cluster_name' in df.columns:
    cleaned_csv.index = df.groupby('cluster')['cluster_name'].first()

```

```
cleaned_csv = cleaned_csv.sort_values('Age_at_Release_numeric')
cleaned_csv
```

Out[25]:

cluster_name	Age_at_Release_numeric	Percent_Days_Employed	Jobs_Per_Year_numeric	Supervision_Programs
High-Risk / Low Employment	30.568253	0.045885	0.167190	
Moderate-Risk / High Employment	31.712495	0.807555	1.182983	
High-Risk / High Employment	32.792531	0.794549	1.242640	
Moderate-Risk / Partial Employment	33.074250	0.493171	0.747794	

In [26]:

```
#Z-score normalization
df_z = df.copy()
df_z[key_features] = (df[key_features] - df[key_features].mean()) / df[key_features].std()


#mean z-scores
cluster_std_centered = df_z.groupby('cluster')[key_features].mean()

if 'cluster_name' in df.columns:
    cluster_std_centered.index = df.groupby('cluster')['cluster_name'].first()

cluster_std_centered = cluster_std_centered.sort_values('Age_at_Release_numeric')

plt.figure(figsize=(10, 6))
sns.heatmap(
    cluster_std_centered,
    annot=True,
    cmap='coolwarm',
    center=0,
    fmt=".2f"
)

plt.title("Z-Score Profiles")
plt.xlabel("Features")
plt.ylabel("Cluster")
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()
```



Data Interpretation

The exploratory data analysis displayed some key distribution across various features; primarily the age, employment, and dependents provided notable information about the potential for recidivism within a 3 year period. The PCA projections provided a preliminary understanding of relationships within a 2-dimensional space. The K-means clustering further discovered four distinct subgroups within the dataset. The clusters were profiled based upon a few key features: Age, percentage employed, jobs per year, supervision risk, dependents, and recidivism within 3 years. The heatmaps of these cluster groups examined the deviations from the mean populations within the data set. The use of the processed z-scores allows for a cleaner comparison of the clusters across the standardized scale, and highlights how significant a deviation of the cluster may be; for example, within the dataset **High-Risk/Low Employment** there is a notable z-score of -0.16 for **Age_at_Release** indicating a slightly younger population. Given the standardized data that was established it can be seen that the cluster with severe negative deviations of employment display higher recidivism rates, whilst high employment experiences a reduced rate.

Conclusion

This data analysis demonstrates that employment is a critical protective factor that can combat recidivism. KMeans clustering, combined with Z-score normalization makes it possible to identify high-risk, low-employment subgroups that could be prioritized for intervention. Given further data such as more detailed arrest records implementing targeted employment programs for these at-risk individuals may reduce recidivism rates and improve

reintegration outcomes. These findings provide actionable insights for policymakers and rehabilitation programs, showing that data-driven clustering can reveal latent patterns to guide interventions.